


(a) Generate 20 data pairs (X, Y) using $y = \sin(2\pi X) + N$


```
In [17]: import numpy as np
import random
random.seed(42)
np.random.seed(42)
X= np.random.uniform(0,1,20)
print (X)
mu,sigma=0,1
N=np.random.normal(mu, sigma, 20)
print(N)
import matplotlib.pyplot as plt
count, bins, ignored = plt.hist(N, 30, density=True)
plt.plot(bins, 1/(sigma * np.sqrt(2 * np.pi)) *
         np.exp( - (bins - mu)**2 / (2 * sigma**2) ),
         linewidth=2, color='r')
plt.show()
```

[0.37454012 0.95071431 0.73199394 0.59865848 0.15601864 0.15599452
0.05808361 0.86617615 0.60111501 0.70807258 0.02058449 0.96990985
0.83244264 0.21233911 0.18182497 0.18340451 0.30424224 0.52475643
0.43194502 0.29122914]
[-1.01283112 0.31424733 -0.90802408 -1.4123037 1.46564877 -0.2257763
0.0675282 -1.42474819 -0.54438272 0.11092259 -1.15099358 0.37569802
-0.60063869 -0.29169375 -0.60170661 1.85227818 -0.01349722 -1.05771093
0.82254491 -1.22084365]



```
In [18]: import math
import matplotlib.pyplot as plt
Y=np.sin(2*np.pi*X)+N
plt.plot(X, Y, 'o', color='red')
```

Out[18]: [



```
In [19]: x_train=X[0:10]
print(x_train)
y_train=Y[0:10]
print(y_train)
x_test=X[10:20]
print(x_test)
y_test=Y[10:20]
print(y_test)
```

```
[0.37454012 0.95071431 0.73199394 0.59865848 0.15601864 0.15599452
 0.05808361 0.86617615 0.60111501 0.70807258]
[-0.3036841  0.00950189 -1.90163109 -1.99324896  2.29630988  0.60480043
 0.42443088 -2.16995176 -1.13782134 -0.8545779 ]
[0.02058449 0.96990985 0.83244264 0.21233911 0.18182497 0.18340451
 0.30424224 0.52475643 0.43194502 0.29122914]
[-1.02201767 0.18776035 -1.46944871 0.68043968 0.30794309 2.76600545
 0.92898565 -1.21263367 1.23723487 -0.25420997]
```

(b) Using root mean square error, find weights of polynomial regression for order is 0, 1, 3, 9 (d) Draw a chart of fit data


```

In [20]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import operator
from operator import itemgetter
from sklearn.metrics import mean_squared_error
from math import sqrt
x1=x_train[:,np.newaxis]
y1=y_train[:,np.newaxis]
poly = PolynomialFeatures(degree = 0)
X_poly = poly.fit_transform(x1)
poly.fit(X_poly, y1)
lin2 = LinearRegression()
lin2.fit(X_poly, y1)
y2=lin2.predict(X_poly)
rms = sqrt(mean_squared_error(y1, y2))
print('RMSE=',rms)
print('Coefficient=',lin2.coef_)
plt.scatter(x1, y1, color='red')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x1,y2), key=sort_axis)
x1, y2 = zip(*sorted_zip)
plt.plot(x1, y2, color='black')
plt.title('When M =0')
plt.xlabel('X')
plt.ylabel('y=sin(2*pi*X)+N')
plt.show()

```

RMSE= 1.3307221036666563

Coefficient= [[0.]]




```

In [21]: poly = PolynomialFeatures(degree = 1)
X_poly = poly.fit_transform(x1)
poly.fit(X_poly, y1)
lin2 = LinearRegression()
lin2.fit(X_poly, y1)
y2=lin2.predict(X_poly)
rms = sqrt(mean_squared_error(y1, y2))
print('RMSE=',rms)
print('Cooefficient=',lin2.coef_)
plt.scatter(x1, y1, color='red')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x1,y2), key=sort_axis)
x1, y2 = zip(*sorted_zip)
plt.plot(x1, y2, color='black')
plt.title('When M =1')
plt.xlabel('X')
plt.ylabel('y=sin(2*pi*X)+N')
plt.show()

```

RMSE= 1.3289257344106773

Cooefficient= [[0. 0.23066742]]




```

In [22]: poly = PolynomialFeatures(degree = 3)
X_poly = poly.fit_transform(x1)
poly.fit(X_poly, y1)
lin2 = LinearRegression()
lin2.fit(X_poly, y1)
y2=lin2.predict(X_poly)
rms = sqrt(mean_squared_error(y1, y2))
print('RMSE=',rms)
print('Coefficient=',lin2.coef_)
plt.scatter(x1, y1, color='red')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x1,y2), key=sort_axis)
x1, y2 = zip(*sorted_zip)
plt.plot(x1, y2, color='black')
plt.title('When M =3')
plt.xlabel('X')
plt.ylabel('y=sin(2*pi*X)+N')
plt.show()

```

RMSE= 1.2185980456718613

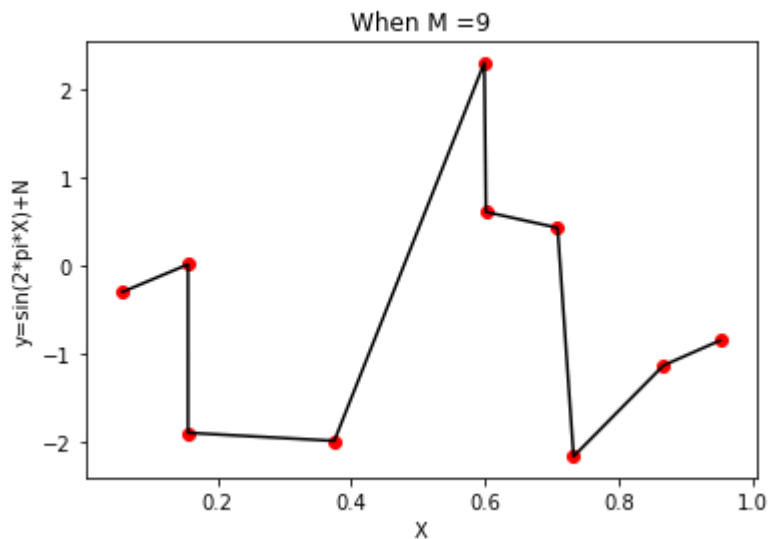
Coefficient= [[0. -11.43208875 33.62650285 -24.62218043]]



```
In [23]: poly = PolynomialFeatures(degree = 9)
X_poly = poly.fit_transform(x1)
poly.fit(X_poly, y1)
lin2 = LinearRegression()
lin2.fit(X_poly, y1)
y2=lin2.predict(X_poly)
rms = sqrt(mean_squared_error(y1, y2))
print('RMSE=',rms)
print('Coefficient=',lin2.coef_)
plt.scatter(x1, y1, color='red')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x1,y2), key=sort_axis)
x1, y2 = zip(*sorted_zip)
plt.plot(x1, y2, color='black')
plt.title('When M =9')
plt.xlabel('X')
plt.ylabel('y=sin(2*pi*X)+N')
plt.show()
```

RMSE= 1.0861921666330498e-07

Coefficient= [[0.00000000e+00 1.76480723e+06 -2.11436975e+07 1.25276252e+08
-4.26092448e+08 8.90250118e+08 -1.16540180e+09 9.33878997e+08
-4.19476059e+08 8.09974015e+07]]



(c) Display weights in table

```
In [24]: from prettytable import PrettyTable
table = PrettyTable()

table.field_names = ["coefficients", "M=0", "M=1", "M=3", "M=9"]

table.add_row(["w0", 0., 0., 0., 0.])
table.add_row(["w1", '', 0.23066742, -11.43208875, 1.76480723e+06])
table.add_row(["w2", '', '', 33.62650285, -2.11436975e+07])
table.add_row(["w3", '', '', -24.62218043, 1.25276252e+08])
table.add_row(["w4", '', '', '', -4.26092448e+08])
table.add_row(["w5", '', '', '', 8.90250118e+08])
table.add_row(["w6", '', '', '', -1.16540180e+09])
table.add_row(["w7", '', '', '', 9.33878997e+08])
table.add_row(["w8", '', '', '', -4.19476059e+08])
table.add_row(["w9", '', '', '', 8.09974015e+07])

print(table)
```

	coefficients	M=0	M=1	M=3	M=9
w0		0.0	0.0	0.0	0.0
w1			0.23066742	-11.43208875	1764807.23
w2				33.62650285	-21143697.5
w3				-24.62218043	125276252.0
w4					-426092448.0
w5					890250118.0
w6					-1165401800.0
w7					933878997.0
w8					-419476059.0
w9					80997401.5

(e) Draw train error vs test error

```

In [25]: xt=x_test[:,np.newaxis]
yt=y_test[:,np.newaxis]
train_error=list()
test_error=list()
degrees=[0,1,2,3,4,5,6,7,8,9]
for i in range(len(degrees)):
    poly = PolynomialFeatures(degree = i)
    X_poly = poly.fit_transform(x1)
    poly.fit(X_poly, y1)
    lin2 = LinearRegression()
    lin2.fit(X_poly, y1)
    y2=lin2.predict(X_poly)
    rms = np.sqrt(mean_squared_error(y1, y2))
    train_error=np.append(train_error,rms)
    rmse_test = np.sqrt(mean_squared_error(yt, lin2.predict(poly.fit_transform
(xt))))
    test_error=np.append(test_error, rmse_test)

print(train_error)
print(test_error)
plt.plot(degrees,train_error, marker='o')
plt.plot(degrees,test_error, marker='o')
plt.ylabel('RMSE')
plt.xlabel('M')


```

```

[1.33072210e+00 1.32892573e+00 1.27891838e+00 1.21859805e+00
 1.10805704e+00 9.61510895e-01 7.40133643e-01 6.05378458e-01
 3.21965094e-01 1.08619217e-07]
[1.41646242e+00 1.45036043e+00 1.43237805e+00 1.83459123e+00
 2.13504028e+00 2.10791197e+00 3.17830978e+00 2.90664081e+01
 1.64069715e+03 7.42911014e+03]

```


Out[25]: Text(0.5, 0, 'M')



(f) Now generate 100 more data and fit 9th order model and draw fit



```
In [26]: random.seed(42)
np.random.seed(42)
A= np.random.uniform(0,1,100)
print (A)
mu,sigma=0,1
B=np.random.normal(mu, sigma, 100)
print(B)
import matplotlib.pyplot as plt
count, bins, ignored = plt.hist(B, 30, density=True)
plt.plot(bins, 1/(sigma * np.sqrt(2 * np.pi)) *
         np.exp( - (bins - mu)**2 / (2 * sigma**2) ),
         linewidth=2, color='r')
plt.show()
```

```
[0.37454012 0.95071431 0.73199394 0.59865848 0.15601864 0.15599452
0.05808361 0.86617615 0.60111501 0.70807258 0.02058449 0.96990985
0.83244264 0.21233911 0.18182497 0.18340451 0.30424224 0.52475643
0.43194502 0.29122914 0.61185289 0.13949386 0.29214465 0.36636184
0.45606998 0.78517596 0.19967378 0.51423444 0.59241457 0.04645041
0.60754485 0.17052412 0.06505159 0.94888554 0.96563203 0.80839735
0.30461377 0.09767211 0.68423303 0.44015249 0.12203823 0.49517691
0.03438852 0.9093204 0.25877998 0.66252228 0.31171108 0.52006802
0.54671028 0.18485446 0.96958463 0.77513282 0.93949894 0.89482735
0.59789998 0.92187424 0.0884925 0.19598286 0.04522729 0.32533033
0.38867729 0.27134903 0.82873751 0.35675333 0.28093451 0.54269608
0.14092422 0.80219698 0.07455064 0.98688694 0.77224477 0.19871568
0.00552212 0.81546143 0.70685734 0.72900717 0.77127035 0.07404465
0.35846573 0.11586906 0.86310343 0.62329813 0.33089802 0.06355835
0.31098232 0.32518332 0.72960618 0.63755747 0.88721274 0.47221493
0.11959425 0.71324479 0.76078505 0.5612772 0.77096718 0.4937956
0.52273283 0.42754102 0.02541913 0.10789143]
[ 0.08704707 -0.29900735 0.09176078 -1.98756891 -0.21967189 0.35711257
1.47789404 -0.51827022 -0.8084936 -0.50175704 0.91540212 0.32875111
-0.5297602 0.51326743 0.09707755 0.96864499 -0.70205309 -0.32766215
-0.39210815 -1.46351495 0.29612028 0.26105527 0.00511346 -0.23458713
-1.41537074 -0.42064532 -0.34271452 -0.80227727 -0.16128571 0.40405086
1.8861859 0.17457781 0.25755039 -0.07444592 -1.91877122 -0.02651388
0.06023021 2.46324211 -0.19236096 0.30154734 -0.03471177 -1.16867804
1.14282281 0.75193303 0.79103195 -0.90938745 1.40279431 -1.40185106
0.58685709 2.19045563 -0.99053633 -0.56629773 0.09965137 -0.50347565
-1.55066343 0.06856297 -1.06230371 0.47359243 -0.91942423 1.54993441
-0.78325329 -0.32206152 0.81351722 -1.23086432 0.22745993 1.30714275
-1.60748323 0.18463386 0.25988279 0.78182287 -1.23695071 -1.32045661
0.52194157 0.29698467 0.25049285 0.34644821 -0.68002472 0.2322537
0.29307247 -0.71435142 1.86577451 0.47383292 -1.1913035 0.65655361
-0.97468167 0.7870846 1.15859558 -0.82068232 0.96337613 0.41278093
0.82206016 1.89679298 -0.24538812 -0.75373616 -0.88951443 -0.81581028
-0.07710171 0.34115197 0.2766908 0.82718325]
```



```
In [27]: C=np.sin(2*np.pi*A)+B  
plt.plot(A, C, 'o', color='red')
```

```
Out[27]: [<matplotlib.lines.Line2D at 0x1e9e3f7d708>]
```




```

In [28]: random.seed(42)
np.random.seed(42)
a1=A[:,np.newaxis]
c1=C[:,np.newaxis]
poly = PolynomialFeatures(degree = 9)
A_poly = poly.fit_transform(a1)
poly.fit(A_poly, c1)
lin2 = LinearRegression()
lin2.fit(A_poly, c1)
c2=lin2.predict(A_poly)
rms = sqrt(mean_squared_error(c1,c2))
print('RMSE=',rms)
print('Coefficient=',lin2.coef_)
plt.scatter(a1, c1, color='red')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(a1,c2), key=sort_axis)
a1, c2 = zip(*sorted_zip)
plt.plot(a1, c2, color='black')
plt.title('When M =9')
plt.xlabel('X')
plt.ylabel('y=sin(2*pi*X)+N')
plt.show()

```

RMSE= 0.8673906703972919




Coefficient= [[0.00000000e+00 3.24189288e+01 -7.13269208e+02 7.17506899e+03
 -3.69590556e+04 1.06904740e+05 -1.81926071e+05 1.80992288e+05
 -9.73972757e+04 2.18912201e+04]]







(g) Now we will regularize using the sum of weights (h) Draw chart for lambda is 1, 1/10, 1/100, 1/1000, 1/10000, 1/100000

```
In [29]: from sklearn.linear_model import Ridge
def plot_lam_graphs(degree,lambda_value):
    polynomial_features= PolynomialFeatures(degree)
    A_Poly = polynomial_features.fit_transform(a1)
    polynomial_features.fit(A_poly,c1)
    model = Ridge(alpha=lambda_value)
    fit = model.fit(A_Poly,c1)
    A_poly_plot = polynomial_features.fit_transform(a1)
    C_pred = model.predict(A_poly_plot)
    plt.plot( a1, C_pred, color='r',label="ln lamda = "+str(lambda_value))
    plt.plot(a1,c1,'o')
    plt.xlabel('x')#xdata ranging from 0 to 1
    plt.ylabel('y')#t is the target value i.e value of y
    plt.title("M="+str(degree))#M no. of coeffecients
    plt.legend()
    plt.show()

lambda_values = np.array([0,1, 0.1, 0.01, 0.001, 0.0001,0.00001])
for i in lambda_values:
    plot_lam_graphs(9,i)
```







(i) Now draw test and train error according to lamda


```

In [30]: a_train=A[0:50]
print(a_train)
c_train=C[0:50]
print(c_train)
a_test=A[50:100]
print(a_test)
c_test=C[50:100]
print(c_test)
a1=a_train[:,np.newaxis]
c1=c_train[:,np.newaxis]
at=a_test[:,np.newaxis]
ct=c_test[:,np.newaxis]

[0.37454012 0.95071431 0.73199394 0.59865848 0.15601864 0.15599452
 0.05808361 0.86617615 0.60111501 0.70807258 0.02058449 0.96990985
 0.83244264 0.21233911 0.18182497 0.18340451 0.30424224 0.52475643
 0.43194502 0.29122914 0.61185289 0.13949386 0.29214465 0.36636184
 0.45606998 0.78517596 0.19967378 0.51423444 0.59241457 0.04645041
 0.60754485 0.17052412 0.06505159 0.94888554 0.96563203 0.80839735
 0.30461377 0.09767211 0.68423303 0.44015249 0.12203823 0.49517691
 0.03438852 0.9093204 0.25877998 0.66252228 0.31171108 0.52006802
 0.54671028 0.18485446]
[ 0.79619409 -0.60375279 -0.90184624 -2.56851417 0.61098923 1.1876893
 1.83479672 -1.26347379 -1.40193222 -1.46725753 1.04437803 0.14081344
-1.39857022 1.48540087 1.00672725 1.88237226 0.24042978 -0.48258489
 0.02258181 -0.49688126 -0.35023069 1.02953751 0.97025763 0.50983789
-1.14284186 -1.39632031 0.60770661 -0.89159569 -0.70985981 0.69178162
 1.26072333 1.05246623 0.65499577 -0.39011493 -2.1330372 -0.95995008
 1.00193024 3.0391318 -1.10819107 0.66878086 0.65911456 -1.13837831
 1.35721495 0.21250582 1.78951068 -1.76210077 2.32855943 -1.5276083
 0.29756298 3.10784668]
[0.96958463 0.77513282 0.93949894 0.89482735 0.59789998 0.92187424
 0.0884925 0.19598286 0.04522729 0.32533033 0.38867729 0.27134903
 0.82873751 0.35675333 0.28093451 0.54269608 0.14092422 0.80219698
 0.07455064 0.98688694 0.77224477 0.19871568 0.00552212 0.81546143
 0.70685734 0.72900717 0.77127035 0.07404465 0.35846573 0.11586906
 0.86310343 0.62329813 0.33089802 0.06355835 0.31098232 0.32518332
 0.72960618 0.63755747 0.88721274 0.47221493 0.11959425 0.71324479
 0.76078505 0.5612772 0.77096718 0.4937956 0.52273283 0.42754102
 0.02541913 0.10789143]
[-1.18048063 -1.55385517 -0.27139852 -1.1172395 -2.12772299 -0.402838
-0.53449822 1.41654712 -0.63906202 2.43999674 -0.13944778 0.66895521
-0.06658335 -0.44750919 1.20863001 1.04208154 -0.83328143 -0.7620665
 0.71135583 0.69952425 -2.22719905 -0.37192468 0.55663109 -0.61961467
-0.71299121 -0.64486535 -1.67110744 0.68088767 1.06969466 -0.04895348
 1.10783635 -0.22567235 -0.31772903 1.04537216 -0.04719498 1.67756761
 0.16679405 -1.58132275 0.31255669 0.58647426 1.50474659 0.92334127
-1.24309298 -1.12931012 -1.88084917 -0.77683674 -0.21945111 0.78085969
 0.43572575 1.45434341]

```

```

In [31]: training=[]
test=[]
lambda_values = np.array([-1,-0.5,-0.3,0,0.00001,0.0001,0.001,0.01,0.1,0.1,0.3
,0.5,1,2])


for i in lambda_values:
    poly_features = PolynomialFeatures(9)
    A_poly = poly_features.fit_transform(a1)
    poly_model = Ridge(alpha=i)
    poly_model.fit(A_poly, c1)
    A_poly_plot = poly_features.fit_transform(a1)
    C_predicted = poly_model.predict(A_poly_plot)
    rmse_train = sqrt(mean_squared_error(c1, C_predicted))
    training.append(rmse_train)
    rmse_test = sqrt(mean_squared_error(ct, poly_model.predict(poly_features.f
it_transform(at))))
    test.append(rmse_test)
print(lambda_values)
print(training)
print(test)
plt.plot(lambda_values,training,'o',linestyle='solid',label='Training')
plt.plot(lambda_values,test,'o',linestyle='solid',label='Test')
plt.xlabel('Lamda')
plt.ylabel('Erms')
plt.legend()
plt.show()

```

```

[-1.e+00 -5.e-01 -3.e-01  0.e+00  1.e-05  1.e-04  1.e-03  1.e-02  1.e-01
 1.e-01  3.e-01  5.e-01  1.e+00  2.e+00]
[1.2385893896410014, 0.8858072216517349, 0.8530988195501283, 0.78696768118439
8, 0.7978915038362975, 0.7983729646604112, 0.8001313381508149, 0.816393812335
3369, 0.8355310591001678, 0.8355310591001678, 0.8437498012252642, 0.852971155
1400033, 0.8789731427498115, 0.9255403600538954]
[1.7179210149514346, 1.2266675316717008, 1.1458094810836168, 1.06664450628627
57, 1.013679662044866, 1.012759315283707, 1.0172652853206696, 1.0419051229633
414, 1.0505210867459562, 1.0505210867459562, 1.0282547175786527, 1.0112995602
427852, 0.9869095410933219, 0.971364834537554]

```



(j)Based on the best test performance, what is your model?

```
In [32]: print('Conclusion: Based on the figure above increasing of lamda is decreasing  
the difference between training error and test error. Based on that we can say  
lamda 2 is the best model from this analysis')
```

Conclusion: Based on the figure above increasing of lamda is decreasing the difference between training error and test error. Based on that we can say lamda 2 is the best model from this analysis